

Linear_Regression_feb14

August 19, 2022

1 Linear Regression

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: data=pd.read_csv("Advertising.csv")
data.head(5)
```

```
[2]:
```

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

```
[3]: data.shape
```

```
[3]: (200, 4)
```

```
[4]: data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   TV           200 non-null    float64
1   radio        200 non-null    float64
2   newspaper    200 non-null    float64
3   sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
[5]: #sns.pairplot(data)# all analysis
```

```
sns.  
↪pairplot(x_vars=["TV","radio","newspaper"],y_vars=["sales"],data=data,size=5)
```


```
/usr/local/lib/python3.7/site-packages/seaborn/axisgrid.py:2076: UserWarning:  
The `size` parameter has been renamed to `height`; please update your code.  
warnings.warn(msg, UserWarning)
```

```
[5]: <seaborn.axisgrid.PairGrid at 0x7f03641f6510>
```




```
[6]: #box plot for all features  
for i in data.columns:  
    sns.boxplot(data[i])  
    plt.show()
```

```
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:  
Pass the following variable as a keyword arg: x. From version 0.12, the only  
valid positional argument will be `data`, and passing other arguments without an  
explicit keyword will result in an error or misinterpretation.  
FutureWarning
```




/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning




```
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:  
Pass the following variable as a keyword arg: x. From version 0.12, the only  
valid positional argument will be `data`, and passing other arguments without an  
explicit keyword will result in an error or misinterpretation.
```

FutureWarning



```
/usr/local/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning:  
Pass the following variable as a keyword arg: x. From version 0.12, the only  
valid positional argument will be `data`, and passing other arguments without an  
explicit keyword will result in an error or misinterpretation.
```

FutureWarning



```
[7]: # outlier Treatment
from scipy import stats
iqr=stats.iqr(data["newspaper"])
q1=stats.scoreatpercentile(data["newspaper"],25)
q3=stats.scoreatpercentile(data["newspaper"],75)
lowerbound=q1 -(1.5*iqr)
upperbound=q3+(1.5*iqr)
print(lowerbound)
print(upperbound)
outliers=data["newspaper"][((data["newspaper"]<lowerbound) | 
    ↳(data["newspaper"]>upperbound))]
print('outliers are',outliers)
```

```
-35.775000000000006
93.625
outliers are 16      114.0
101      100.9
Name: newspaper, dtype: float64
```

```
[8]: median = np.median(data["newspaper"])# Replace with median

for i in outliers:
    data["newspaper"] = np.where(data["newspaper"]==i, median,
    ↳data["newspaper"])
```

```
[9]: # Simple Linear Regression

# consider one features

# Tv as feature and sale as a label
```

```
[10]: # split input and output features

X=data.loc[:,["TV"]]
y=data.loc[:,["sales"]]
```

```
[11]: X.head()
```

```
[11]:      TV
0  230.1
1   44.5
2   17.2
3  151.5
4  180.8
```

2 split the data for training and Testing

```
[31]: # Scikit learn----- it contain all machine learning realted libraries

from sklearn.model_selection import train_test_split

xtrain,xtest,ytrain,ytest=train_test_split(X,y,test_size=0.3,random_state=10)
```

```
[32]: print(xtrain.shape)
print(xtest.shape)
print(ytrain.shape)
print(ytest.shape)
```

```
(140, 1)
(60, 1)
(140, 1)
(60, 1)
```

```
[33]: # Build the Model
#process to import the linear regression model
from sklearn.linear_model import LinearRegression

# instantiate the object

lr=LinearRegression()
```

```
[34]: # fit the data or train the data  
lr.fit(xtrain,ytrain)
```

```
[34]: LinearRegression()
```

```
[35]: #  $y=mx+c$   
lr.intercept_
```

```
[35]: array([7.30826074])
```

```
[36]: lr.coef_
```

```
[36]: array([[0.04474675]])
```

```
[37]: # Sales = 0.048 * [TV] + 6.917
```

```
[38]: #training performance  
ypredtrain=lr.predict(xtrain)
```

```
[ ]: #ypredtrain
```

```
[39]: ypredtrain.shape
```

```
[39]: (140, 1)
```

```
[40]: ypredtrain
```

```
[40]: array([[ 8.17634775],  
          [17.41207755],  
          [14.6154055 ],  
          [16.56188925],  
          [14.77201913],  
          [17.93113989],  
          [17.8013743 ],  
          [10.80745682],  
          [13.54148342],  
          [ 8.71778346],  
          [13.63097693],  
          [ 9.23684579],  
          [11.52340487],  
          [11.14305747],  
          [13.35802174],  
          [16.15021912],  
          [10.36893864],  
          [ 8.10922762],  
          [10.72691267],
```

[15.22843601],
[10.72691267],
[18.03853209],
[11.65764513],
[9.54559839],
[17.47472301],
[19.07218209],
[15.47901783],
[13.47883797],
[19.89999702],
[14.00684966],
[11.80083474],
[16.44107301],
[10.67769124],
[16.86169249],
[18.66498664],
[10.4673815],
[14.08739381],
[17.52394444],
[18.19067105],
[20.10583208],
[10.65084319],
[8.14502502],
[13.55043278],
[15.85936523],
[15.14341718],
[7.68413347],
[15.58193536],
[14.71832303],
[10.68664059],
[20.4459074],
[9.97516722],
[19.17509962],
[12.49888408],
[15.02707562],
[14.92415809],
[11.9887711],
[16.94671132],
[8.37323346],
[15.17026523],
[16.9153886],
[7.89891788],
[8.42692957],
[15.95333341],
[18.05195612],
[9.88119904],
[17.16597041],

[9.3084406],
[7.49172243],
[9.00863735],
[17.96693729],
[11.590525],
[7.89444321],
[16.24866198],
[8.95941593],
[11.61289837],
[13.58623018],
[15.39847367],
[17.14359704],
[8.43140424],
[16.15021912],
[18.42335417],
[18.44125287],
[9.46952891],
[14.6691016],
[14.31560225],
[12.22145421],
[12.93292759],
[8.07790489],
[15.22843601],
[19.55544702],
[16.85721782],
[9.82302826],
[12.11406201],
[12.72261785],
[19.86867429],
[14.84361394],
[18.53522105],
[10.40473604],
[18.05195612],
[19.2511691],
[13.40276849],
[12.36016915],
[10.39578669],
[10.59267241],
[12.22145421],
[7.83179775],
[20.41458467],
[10.72243799],
[17.04962885],
[15.97570678],
[7.65728542],
[14.97337952],
[11.67106915],


```
[12.70024447],
[17.08542626],
[20.17742688],
[19.19299832],
[18.01615872],
[17.3315334 ],
[11.35336721],
[11.25939903],
[19.06323274],
[15.58193536],
[17.25993859],
[10.34209059],
[10.66874189],
[12.55258019],
[16.36947821],
[17.60448859],
[13.09849057],
[ 7.69308282],
[ 9.0175867 ],
[14.01132433],
[11.50998084],
[12.81658603],
[16.68718016],
[13.17456005],
[16.05177626],
[11.2101776 ],
[16.24866198]])
```

```
[41]: from sklearn.metrics import mean_squared_error
print("The Rms value is training data")
print("-----")
print(np.sqrt(mean_squared_error(ytrain,ypredtrain)))
```

```
The Rms value is training data
-----
3.0681028951223004
```

```
[42]: plt.scatter(xtrain,ytrain)
plt.plot(xtrain,ypredtrain,color="red")
```

```
[42]: [<matplotlib.lines.Line2D at 0x7f031403e310>]
```



```
[43]: # testing the data
ypredtest=lr.predict(xtest)
```

```
[44]: ypredtest.shape
```

```
[44]: (60, 1)
```

```
[45]: from sklearn.metrics import mean_squared_error
print("The Rms value is test data")
print("-----")
print(np.sqrt(mean_squared_error(ytest,ypredtest)))
```


```
The Rms value is test data
```

```
-----
```

```
3.6508293924109783
```

```
[46]: plt.scatter(xtest,ytest)
plt.plot(xtest,ypredtest,color="red")
```

```
[46]: [<matplotlib.lines.Line2D at 0x7f0313fb6310>]
```



```
[47]: from sklearn.metrics import r2_score
      #r2score close to 1 is a good value
      print(r2_score(ytest,ypredtest))
```

0.6447879078216139

```
[ ]: # End of Simple Linear Regression
```

3 Multiple Linear Regression

```
[48]: data.head(2)
```

```
[48]:      TV  radio  newspaper  sales
0  230.1   37.8     69.2    22.1
1   44.5   39.3     45.1    10.4
```

```
[49]: # splitted i/p and o/p

xnew=data.drop(["sales"],axis=1) #drop the column
ynew=data["sales"]
```

```
[ ]: xnew.head(3)
```

```
[50]: from sklearn.model_selection import train_test_split

xtrain,xtest,ytrain,ytest=train_test_split(xnew,ynew,test_size=0.
↳3,random_state=10)
```

```
[51]: print(xtrain.shape)
print(xtest.shape)
```

```
(140, 3)
(60, 3)
```

```
[52]: # build a mlr model
from sklearn.linear_model import LinearRegression

mlr=LinearRegression()

mlr.fit(xtrain,ytrain)
```

```
[52]: LinearRegression()
```

```
[53]: print("intercept value",mlr.intercept_)
print("coefficent value",mlr.coef_)
```

```
intercept value 3.392597098985595
coefficent value [ 0.04304513  0.19298948 -0.00301967]
```

```
[ ]: # sales= 0.043 [TV]+[0.193][radio]-[0.0038]newspaper+ 3.41
```

```
[54]: # check the train performance
```

```
ypredtraining_mlr=mlr.predict(xtrain)
```

```
[55]: df=pd.DataFrame({"Actual":ytrain,"predicted":ypredtraining_mlr})
df
```

```
[55]:
```

	Actual	predicted
119	6.6	7.248166
110	13.4	14.524089
95	16.9	16.360594
167	12.2	13.239294
193	19.6	18.667212
..
113	15.9	16.358129
64	18.0	17.208495
15	22.4	20.849473
125	10.6	9.345199
9	10.6	12.430770

[140 rows x 2 columns]

```
[56]: from sklearn.metrics import mean_squared_error
print("The Rms value is training data")
print("-----")
print(np.sqrt(mean_squared_error(ytrain,ypredtraining_mlr)))
```

The Rms value is training data

1.3332818533291317

```
[57]: from sklearn.metrics import r2_score
#r2score close to 1 is a good value
print(r2_score(ytrain,ypredtraining_mlr))
```

0.920909492761702

```
[58]: # Test performace

ypredtest_mlr=mlr.predict(xtest)
```

```
[59]: pd.DataFrame({"Actual":ytest,"predicted":ypredtest_mlr}).head()
```

```
[59]:
```

	Actual	predicted
59	18.4	18.127313
5	7.2	12.977800
20	18.0	17.978211
198	25.5	23.505852
52	22.6	20.635645

```
[60]: from sklearn.metrics import mean_squared_error
print("The Rms value is test data")
print("-----")
print(np.sqrt(mean_squared_error(ytest,ypredtest_mlr)))
```

The Rms value is test data

2.314415304470104

```
[61]: print(r2_score(ytest,ypredtest_mlr))
```

0.8572463405728885

```
[ ]: # End of Linear regression
```

```
[ ]:
```